

Nearest neighbor search in high dimensional spaces

Ilya Fedorov

March 2020

- Why do we need to optimize KNN?
- Data structures
- Complexity analysis
- Difficulties in high dimensional spaces
- Approximate NN search: Locality Sensitive Hashing
- Approximate NN search: FlyHash
- Approximate NN search: Hierarchical Navigable Small World

Why do we need to optimize KNN?

- Very large datasets
- Very frequent queries
- Duplicates search
- It seems that brute force algorithm doesn't use information gained from calculating previous distances: if $d(a, b)$ is big and $d(a, c)$ is small does it mean that $d(c, b)$ is big?

- KD-Tree
- Ball-tree
- Ball*-tree
- R-Tree
- etc...


algorithm : {'auto', 'ball_tree', 'kd_tree', 'brute'}, optional

Algorithm used to compute the nearest neighbors:

- 'ball_tree' will use `BallTree`
- 'kd_tree' will use `KDTree`
- 'brute' will use a brute-force search.
- 'auto' will attempt to decide the most appropriate algorithm based on the values passed to `fit` method.

Note: fitting on sparse input will override the setting of this parameter, using brute force.


Data structures




- Tree construction - by *build_node*($\{(x_1, y_1), \dots (x_N, y_N)\}$):

```
build_node( $\Omega$ ): # omega-objects in the node
    if  $|\Omega| < n_{min}$ :
        node.objects =  $\Omega$ 
    else:
        find feature with maximal spread in  $\Omega$ :
             $x^i = \arg \max_{x^i} \sigma(x^i)$ 
        find median  $\Omega$ :  $\mu = \text{median}\{x^i\}$  # yields balanced tree
        node.feature = i
        node.threshold =  $\mu$ 
        node.left child =
            build_node( $\{x_k \in \Omega : x_k^i < \mu\}$ )
        node.right child =
            build_node( $\{x_k \in \Omega : x_k^i \geq \mu\}$ )
    return node
```


Data structures




A ball-tree: level 4



Data structures



(a) Ball-tree




(b) Ball*-tree

Fig. 2: Comparison of the splitting algorithms in ball-tree and ball*-tree



Dolatshah, Mohamad Hadian, Ali Minaei, Behrouz. (2015). Ball*-tree: Efficient spatial indexing for constrained nearest-neighbor search in metric spaces.

R-tree



Complexity analysis

Fit + Predict time, 10000 objects



KD-Tree timings, 10000 objects

